

Leilani H. Gilpin

Engineering 2 Office 347B
1156 High Street
Santa Cruz, CA 95060
☎ +1 (415) 937 1806
✉ lgilpin@ucsc.edu
🌐 lgilpin.com
🐦 [leilanigilpin](https://twitter.com/leilanigilpin)

Last updated: November 6, 2024

Research Interests

The theories and methodologies towards monitoring, designing, and augmenting machines that can **explain** themselves for diagnosis, accountability, and liability.

Employment

- 2021-present **UC Santa Cruz**, Assistant Professor of Computer Science and Engineering (CSE)
Faculty affiliate in the Science and Justice Research Center (SJRC).
- 2020-2021 **Sony AI**, Research Scientist II
- 2013-2015 **Palo Alto Research Center (PARC)**, Member of Technical Staff

Education

- 2015–2020 **Ph.D., Electrical Eng. and Computer Science, MIT**
Dissertation: Anomaly Detection through Explanations
Advisors: Gerald Jay Sussman and Lalana S. Kagal
- 2011–2013 **M.S., Computational and Mathematical Engineering, Stanford University**
- 2006–2011 **B.S., Computer Science and B.S., Mathematics, UC San Diego (UCSD)**
Highest Honors in Computer Science, Honors in Mathematics, Music Minor

Honors and Awards

- 2024 Air Force Office of Scientific Research Young Investigator Program (AFOSR YIP)
- 2022 AAI New Faculty Highlight
- 2020 Rising Stars in EECS
- 2020 ACM FAccT Travel Award
- 2018 Nokia Bell Labs Prize Finalist
Finalist for prize that recognizes research that “changes the game” in the field of information and communications technologies by a factor of 10.
- 2018 AAI Doctoral Consortium Travel Award
- 2017 Nokia Bell Labs Prize Semi-finalist

2016	USENIX Security Student Travel Award
2016-2020	MIT University Center for Exemplary Mentoring (UCEM) Sloan Scholar
2015	MIT ODGE Diversity Fellowship
2011-2013	National Science Foundation (NSF) Graduate Research Fellowship
2013	Stanford SSB Health IT Competition 1st Place
2011	Stanford School of Engineering Fellowship
2011	Yahoo! HackU All Stars Finalist
2011	Yahoo! HackU First Place
2011	Yahoo! Excellence Award
2010	CRA Outstanding Undergraduate Researcher Honorable Mention
2009-present	Member of Tau Beta Pi and Eta Kappa Nu
2010	Tau Beta Pi Scholarship
2009	Gary C. Reynolds Memorial Scholarship
2009	BAE Scholarship Finalist

Publications

According to Google scholar: 3,615 citations, h-index: 10 , i10-index: 11.

Refereed Journal Publications

- [1] Adam Amos-Binks, Dustin Dannenhauer, and Leilani H Gilpin. "The anticipatory paradigm". In: *AI magazine* 44.2 (2023-06). ISSN: 0738-4602.
- [2] Sharan Subramanian and Leilani Gilpin. "Convolutional Neural Network Model for Diabetic Retinopathy Feature Extraction and Classification". In: *Journal of Student Research* 13.1 (Feb. 2024). DOI: 10.47611/jsrhs.v13i1.6093. URL: <https://www.jsr.org/hs/index.php/path/article/view/6093>.
- [3] Ryan Jenkins, Kristian Hammond, Sarah Loehr, and Leilani Gilpin. "Separating Facts and Evaluation: Motivation, Account, and Learnings from a Novel Approach to Evaluating the Human Impacts of Machine Learning". In: *AI and Society*. 2022.
- [4] Peter R Wurman, Samuel Barrett, Kenta Kawamoto, James MacGlashan, Kaushik Subramanian, Thomas J Walsh, Roberto Capobianco, Alisa Devlic, Franziska Eckert, Florian Fuchs, Leilani Gilpin, et al. "Outracing champion Gran Turismo drivers with deep reinforcement learning". In: *Nature* 602.7896 (2022), pp. 223–228.
- [5] Michal Araszkievicz, Ilaria Angela Amantea, Saurabh Chakravarty, Robert van Doesburg, Maria Dymitruk, Marie Garin, Leilani Gilpin, Daphne Odekerken, and Seyedeh Sajedeh Salehi. "ICAIL Doctoral Consortium, Montreal 2019". In: *Artif. Intell. Law* 28.2 (2020), pp. 267–280. DOI: 10.1007/s10506-020-09267-z. URL: <https://doi.org/10.1007/s10506-020-09267-z>.
- [6] Ioana Baldini, Clark Barrett, Antonio Chella, Carlos Cinelli, David Gamez, Leilani Gilpin, Knut Hinkelmann, Dylan Holmes, Takashi Kido, Murat Kocaoglu, et al. "Reports of the AAAI 2019 Spring Symposium Series". In: *AI Magazine* 40.3 (2019), pp. 59–66.

- [7] Leilani H. Gilpin, Jamie C. Macbeth, and Evelyn Florentine. "Monitoring Scene Understanders with Conceptual Primitive Decomposition and Commonsense Knowledge". In: *Advances in Cognitive Systems* 6 (2018).
- [8] Ayesha Bose, Leilani Gilpin, Jamin Agosti, and Quinn Dang. "The Veicl Act: Safety and Security for Modern Vehicles". In: *Willamette L. Rev.* 53 (2016), p. 137.
- [9] Juan Liu, Eric Bier, Aaron Wilson, John Alexis Guerra-Gomez, Tomonori Honda, Kumar Sricharan, Leilani Gilpin, and Daniel Davies. "Graph analysis for detecting fraud, waste, and abuse in healthcare data". In: *AI Magazine* 37.2 (2016), pp. 33–46.

Peer Reviewed Conference Publications

- [10] Li Liu, Diji Yang, Sijia Zhong, Kalyana Suma Sree Tholeti, Lei Ding, Yi Zhang, and Leilani H. Gilpin. "Right this way: Can VLMs Guide Us to See More to Answer Questions?" In: *To appear in NeurIPS 2024*. 2024.
- [11] Priyesh Vakharia, Abigail Kufeldt, Max Meyers, Ian Lane, and Leilani H Gilpin. "ProSLM: A Prolog Synergized Language Model for explainable Domain Specific Knowledge Based Question Answering". In: *International Conference on Neural-Symbolic Learning and Reasoning*. Springer. 2024, pp. 291–304.
- [12] Rui-Jie Zhu, Ziqing Wang, Leilani Gilpin, and Jason K. Eshraghian. "Autonomous Driving with Spiking Neural Networks". In: *To appear in NeurIPS 2024*. 2024. URL: <https://arxiv.org/abs/2405.19687>.
- [13] Diego Ortiz Barbosa, Leilani H. Gilpin, and Alvaro A. Cardenas. "Semi-Automated Synthesis of Driving Rules". In: *ISOC Symposium on Vehicle Security and Privacy (VehicleSec 2023)* (2023).
- [14] Biagio La Rosa, Leilani H. Gilpin, and Roberto Capobianco. "Towards a fuller understanding of neurons with Clustered Compositional Explanations". In: *Thirty-seventh Conference on Neural Information Processing Systems*. 2023. URL: <https://openreview.net/forum?id=51PLYhMFwz>.
- [15] Adam Amos-Binks, Dustin Dannenhauer, and Leilani H Gilpin. "Anticipatory Thinking Challenges in Open Worlds: Risk Management". In: *AAAI Spring Symposium on Designing Artificial Intelligence for Open Worlds*. 2022.
- [16] Gregory Falco and Leilani H Gilpin. "A stress testing framework for autonomous system verification and validation (v&v)". In: *2021 IEEE International Conference on Autonomous Systems (ICAS)*. IEEE. 2021, pp. 1–5.
- [17] Leilani H. Gilpin, Vishnu Penubarthi, and Lalana Kagal. "Explaining Multimodal Errors in Autonomous Vehicles". In: *2021 IEEE 8th International Conference on Data Science and Advanced Analytics (DSAA)*. 2021, pp. 1–10. DOI: 10.1109/DSAA53316.2021.9564178.
- [18] Leilani H. Gilpin. "System-wide Monitoring for Anomaly Detection". In: *Advances in Cognitive Systems* (2020).
- [19] Jason R. Wilson, Leilani H. Gilpin, and Irina Rabkina. "A Knowledge Driven Approach to Adaptive Assistance Using Preference Reasoning and Explanation". In: *arXiv preprint arXiv:2012.02904* (2020).

- [20] Leilani H. Gilpin. "Explaining Possible Futures for Robust Autonomous Decision-Making". In: *Proceedings of the AAAI Fall Symposium on Anticipatory Thinking* (2019).
- [21] Leilani H. Gilpin, David Bau, Ben Z Yuan, Ayesha Bajwa, Michael Specter, and Lalana Kagal. "Explaining explanations: An overview of interpretability of machine learning". In: *2018 IEEE 5th International Conference on data science and advanced analytics (DSAA)*. IEEE. 2018, pp. 80–89.
- [22] Leilani H. Gilpin, Danielle M. Olson, and Tarfah Alrashed. "Perception of Speaker Personality Traits Using Speech Signals". In: *Extended Abstracts of the 2018 CHI Conference on Human Factors in Computing Systems*. ACM. 2018, LBW514.
- [23] Leilani H. Gilpin, Cagri Zaman, Danielle Olson, and Ben Z Yuan. "Reasonable perception: Connecting vision and language systems for validating scene descriptions". In: *Companion of the 2018 ACM/IEEE International Conference on Human-Robot Interaction*. ACM. 2018, pp. 115–116.
- [24] Leilani H. Gilpin and Ben Ze Yuan. "Getting Up to Speed on Vehicle Intelligence". In: *AAAI Spring Symposium Series*. 2017.
- [25] Leilani Gilpin, Laurent Ciarletta, Yannick Presse, Vincent Chevrier, and Virginie Galtier. "Co-simulation solutions using AA4MM-FMI applied to smart space heating models". In: *Seventh International Conference on Simulation Tools and Techniques*. 2014.

Peer Reviewed Short Conference Publications and Abstracts

- [26] Jonathan Wellington Morris, Vishrut Shah, Alex Besanceney, Daksh Shah, and Leilani H. Gilpin. *Slug Mobile: Test-Bench for RL Testing*. 2024. arXiv: 2409.10532 [cs.R0]. URL: <https://arxiv.org/abs/2409.10532>.
- [27] Leilani H Gilpin. "Accountability layers: explaining complex system failures by parts". In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 37. 13. 2023, pp. 15439–15439.
- [28] Leilani H. Gilpin. "Anticipatory Thinking: A Testing and Representation Challenge for Self-Driving Cars". In: *2021 55th Annual Conference on Information Sciences and Systems (CISS)*. 2021, pp. 1–2. DOI: 10.1109/CISS50987.2021.9400212.
- [29] Leilani H. Gilpin and Lalana Kagal. "An Adaptable Self-Monitoring Framework for Opaque Machines". In: *Proceedings of the 18th International Conference on Autonomous Agents and MultiAgent Systems*. International Foundation for Autonomous Agents and Multiagent Systems. 2019, pp. 1982–1984.
- [30] Leilani H. Gilpin. "Reasonableness Monitors". In: *Thirty-Second AAAI Conference on Artificial Intelligence*. 2018.

Peer Reviewed Workshop Publications

- [31] Shengjie Xu and Leilani H. Gilpin. "DANGER: A Framework for Generating Dangerous Scenes for Testing Robustness of Machine Learning". In: *Neurips Workshop on Progress and Challenges in Building Trustworthy Embodied AI* (2022).

- [32] Leilani H. Gilpin. "Reconciling System-wide Errors with Symbolic Explanations". In: *Proceedings of the IJCAI Workshop on AI for Anomaly Detection* (2020).
- [33] Leilani H. Gilpin. "Monitoring Opaque Learning Systems". In: *ICLR 2019 Debugging ML Models Workshop* (2019).
- [34] Leilani H. Gilpin, Tianye Chen, and Lalana Kagal. "Learning From Explanations for Robust Autonomous Driving." In: *ICML Workshop on AI for Autonomous Driving*. 2019.
- [35] Leilani H. Gilpin, Cecilia Testart, Nathaniel Fruchte, and Julius Adebayo. "Explaining Explanations to Society". In: *arXiv preprint arXiv:1901.06560* (2018).

Other Publications and Pre-Prints

- [36] Shiyuan Huang, Siddarth Mamidanna, Shreedhar Jangam, Yilun Zhou, and Leilani H. Gilpin. *Can Large Language Models Explain Themselves? A Study of LLM-Generated Self-Explanations*. 2023. arXiv: 2310.11207 [cs.CL].
- [37] Shreyan Mitra and Leilani Gilpin. *A novel post-hoc explanation comparison metric and applications*. 2023. arXiv: 2311.10811 [cs.LG].
- [38] Shreyan Mitra and Leilani Gilpin. *The XAISuite framework and the implications of explanatory system dissonance*. 2023. arXiv: 2304.08499 [cs.LG].
- [39] Sharan Subramanian and Leilani H. Gilpin. *Convolutional Neural Network Model for Diabetic Retinopathy Feature Extraction and Classification*. 2023. arXiv: 2310.10806 [eess.IV].
- [40] Karianne Bergen and Leilani Gilpin. *Negative News No More: Classifying News Article Headlines*. Tech. rep. 11. 2012, p. 2017.
- [41] Leilani Gilpin. *The Impact of Topology and Communication Models on Connectivity in Networks*. 2011.
- [42] Leilani Hendrina Gilpin. "Anomaly detection through explanations". PhD thesis. Massachusetts Institute of Technology, 2020.

Invited Talks and Panels

NeuroSymbolic Approaches for Trustworthy, Explainable Complex Systems

Talks Alan Turing Institute NeSy AI group (Remote) 2024

Responsible Explainability for Complex Systems

adeptID Lunch & Learn 2024

Explaining Self-Driving Cars for Accountable Autonomy

NeurIPS Workshop on XAI in Action 2023

Explaining and Generating Dangerous Scenarios for Testing AVs

Workshop on Scenario Generation for Testing AVs (SceGen) 2023

Kansas State University 2023

Accountability Layers

Stanford University: Computing & Society Seminar 2024

	Brown University	2023
	UCSC Applied Mathematics Seminar	2023
	ML Fairness at Illinois University	2023
	University of Birmingham (joint with R. Marinescu)	2022
	University College London	2022
	Explaining Errors in Autonomous Driving	
	JHU Department of Civil & Systems Engineering	2021
	Anomaly Detection through Explanations	
	55th Annual Conference on Information Sciences and Systems	2021
	FUZZ-IEEE	2021
	Cornell Computing and Information Science	2021
	Utah School of Computing	2021
	USC Computer Science	2021
	UC Santa Cruz CSE	2021
	Institute of Assured Autonomy at JHU	2020
	31st International Workshop on Principles of Diagnosis (DX)	2020
	USC - Information Science Institute	2020
	ICML Workshop on Monitoring ML Systems	2020
	Sony AI	2020
	Northeastern University - Experiential AI	2020
	Northwestern University - Computer Science	2020
	Salesforce Research (canceled due to Covid-19)	2020
	Rochester Data Science Consortium (postponed due to Covid-19)	2020
	Idexx (postponed due to Covid-19)	2020
	Explaining Explanations in AI	
	Statistics Canada (STATCAN)	2021
	TU Berlin	2020
	U. of Cambridge Distributed Info. and Automation Lab (DIAL) XAI Day	2020
	Stanford University - Knowledge Graphs Seminar	2020
	The Car Can Explain!	
	UC San Diego - Halicioglu Data Science Institute	2019
	CSAIL-Toyota Meeting	2018
	MIT Museum	2018
	Columbia Law School: Software Freedom Law Center	2018
	CSAIL-Toyota Meeting	2017
Panelist	UCSC Open Source Symposium on AI and Open Source	2024

	ICLR Workshop on AI for Earth and Space Sciences	2022
	ESIP Winter Meeting visionary panel on knowledge graphs	2021
	Northwestern ML Impact Initiative: Intelligibility, Fairness, Transparency	2021
	Columbia Law School: Software Freedom Law Center	2019
Invite-only	Bay Area Alignment Workshop	2024
	Dagstuhl Seminar	2024
	<i>Trustworthiness and Responsibility in AI - Causality, Learning, and Verification</i>	
	8th Annual CHAI Workshop (Asilomar)	2024
Outreach	Rotatary Club of Salinas	2024
	BayHacks	2024
	UCSC Slugs and Steins	2023
	UCSC AI Club: XAI Talk	2022
	St. Olaf Hackathon	2023
	Baskin Day	2023

Grants

- PI **Air Force Office of Scientific Research Young Investor Program (AFOSR YIP).** *Frame-Based Monitoring to Detect and Explain Multimodal Autonomous System Errors.* March 1, 2024-February 28, 2027. Award Amount: \$450,000.00.
- Underwriters Laboratory (subaward under Northwestern University).** *Adversarial Examples to Test Explanation Robustness.* March 1, 2022-April 1, 2023. Award Amount: \$125,477.00.
- UCSC Office of Research Seed Funding for Early Stage Initiatives.** *Robustifying Machine Learning for Safe and Secure Autonomous Vehicles.* March 15, 2022 - March 31, 2023. Alvaro Cardenas (Co-PI), Daniel Fremont (Co-PI), and Cihang Xie (Co-PI). Award Amount: \$40,000.00.
- Underwriters Laboratory (subaward under Northwestern University).** *Explaining Hallucinations in LLMS with Commonsense Monitoring.* March 1, 2024-April 1, 2025. Award Amount: \$150,000.00.
- Co-PI **Department of Transportation (DoT) National Center for Transportation Cybersecurity and Resiliency.** *Reducing Transportation Cybersecurity Risks.* 2022-2026. UC Santa Cruz Team: Alvaro Cardenas (Co-PI), Daniel Fremont (Co-PI), and Cihang Xie (Co-PI).
- Learning Lab's Grand Challenge: Building Critical Mass for Data Science.** *Building Data Science Communities for Improving Student Success.* 2023-2026. Team: Pedro Morales-Almazan(PI), Marcella Gomez (Co-PI), Judith Canner (Co-PI), Leilani Gilpin (Co-PI), Sangwon Hyun (Co-PI). Award Amount: \$345,000.00

Gifts

Siemens CITRIS Research Gift 2023-2024. Gift Amount: \$50,000

UCSC Committee on Research (COR) Faculty Allowance (CFA). 2022, 2023.
Gift Amount: \$2,000 per year.

Tutorials

AAAI '23	Trustworthy and Responsible AI [tutorial website]	2023
	<i>Co-lead with Yilun Zhou, Jieyu Zhao, Harsha Nori, and Besmira Nushi.</i>	
US2TS '22	Knowledge-based commonsense reasoning and explainability	2022
	<i>Co-lead with Filip Ilievski of USC ISI.</i>	

Selected Press

- o Awe and trepidation as AI comes for smartphones: [\[Yahoo! News\]](#)
- o The Tricky Ethics of AI in the lab: [\[Chemical and Engineering News \(CEN\)\]](#)
- o Celebrating Baskin Engineers on International Women in Engineering Day: [\[UCSC News\]](#)
- o New data science project seeks to span disciplines and create community: [\[UCSC News\]](#).
- o UC Santa Cruz engineers join major transportation cybersecurity project: [\[UCSC News\]](#).
- o Santa Cruz County's potential vulnerability to a cyber attack: [\[KION 46 article\]](#)
- o Researchers Develop DANGER Framework for Stress Testing. [\[Research Spotlight\]](#)
- o Robustifying AVs Project Chosen for Seed Funding. [\[UCSC News\]](#) .
- o Mob.ly App Makes Driving Safer by Changing How Drivers Navigate. [\[AI Pulse Report\]](#) .
- o MIT CSAIL Student Spotlight. [\[Student Profile\]](#)
- o MIT student lead AI and Ethics Reading Group. [\[MIT News\]](#).
- o MIT Internet Policy Research Initiative (IPRI) [\[Student profile\]](#).

Teaching

Lead Instructor

UCSC	Advanced ML (CSE 290C) Graduate	Spring 2024
	Artificial Intelligence (CSE 240) Graduate	Winter 2023, Winter 2024, Fall 2024
	Responsible Data Science (CSE 246)	Fall 2022
	Artificial Intelligence (CSE 140)	Winter 2021, Spring 2022/2023, Fall 2023
MIT	Artificial Intelligence and Global Risks	IAP 2018

Developed, taught, managed a new course on the risks of AI from a global perspective. [course webpage].

Stanford SMASH Institute - Calculus Summer 2015
Planned and lead weekly lectures to teach a semester-long calculus class over the summer.

Guest Lectures

UCSC Crown College plenary 2024
Navigating the Ethical Landscape of AI a panel with Sri Rao
Artificial Intelligence: CSE140 Spring 2024
Delivered two lectures on Adversarial Search.
Advanced Computer Security: CSE233 Winter 2024
Delivered two lectures on (deep) Reinforcement Learning.
Machine Learning (Graduate): CSE 242 Fall 2022
Machine Learning: CSE 142 Fall 2022, Fall 2023
Projects in Artificial Intelligence: CSE 247 Spring 2022
Research & Teaching in CSE: CSE 200 Winter 2022, Fall 2022, Fall 2023
UConn Projects in Artificial Intelligence: 5095 - IETs Fall 2022
USC DSCI-558 class on Common Sense and Explainability Fall 2022
Emmanuel Practical Machine Learning (Invited talk on GANs and XAI) Spring 2022/2023
Northwestern CS 496: AI Perspectives (Invited talk on XAI) Spring 2021/2022/2023
Stanford CS 520: Knowledge Graphs (Invited talk on XAI) Spring 2020
MIT 6.905/6.945: Large-scale Symbolic Systems Spring 2019
6.S978: Privacy Legislation in Practice: Law and Technology Spring 2017

Teaching Assistant

MIT 6.905/6.945: Large-scale Symbolic Systems Spring 2019
Stanford CS 348A: Geometric Modeling (PhD Level Course) Spring 2013
UCSD COGS 5A (beginning java)
CSE 8A/8B (beginning java)
CSE 5A (beginning C)
CSE 21 (discrete mathematics)
CSE 100 (Advanced Data Structures)
CSE 101 (Algorithms)

Mentoring

PhD students (primary supervisor)

	Oliver Chang	2022-present
	Li Liu	2023-present
	Jack Fox Keen	2023-present
	Peiyu (Olivia) Wang	2023-present
	Shiyuan Huang	2024-present
	Priyesh Vakharia	2024-present
	Creighton Glasscock	2024-present (incoming)
	Postdoctoral Researchers (primary supervisor)	
	Biagio La Rosa	2024-present
	Visiting PhD students (primary supervisor)	
Sapienza U.	Biagio La Rosa	2022-2023
	PhD student committee service (other supervisor)	
Def. Committee	Connor Priyor	2024
	Eliana Stefani	2024
	Jayanth Yetukuri	2024
	Zhaowei Zhu	2023
Adv. Committee	Héctor Carrión	2024
	Sebastián Castro	2024
	Zonglin Di	2024
	Changmao Li	2023
	Michael Briden	2022
	Anuj Kaul	2022
	Master's Students	
Thesis Advisor	Sijia Zhong	2023-present
	Sanjana Sudhakar Patil	2023-2024
	<i>Co-advised with Kate Ringland, Prof. of Computational Media</i>	
	Shengjie Xu	2021-2023
	Sainaya Brid	2022-2023
	Sanjana Sudhakar Patil	2023-2024
	<i>Co-advised with Kate Ringland, Prof. of Computational Media</i>	
Thesis Reader	Ex Taranenko	2024
Project Advisor	Batuhan Salih	2024-present
	Aaja Ouellette	2024-present
	Ziyuan Wang	2024-present
	Rithesh Kumar	2024-present
	Akashaleena Sarkar	2024-present

Zhonghui Li	2024-present
Vinh Le	2024-present
Pratik Doshi	2024-present
Kajal Patil	2023-present
Karthik Bhatt	2024-present
William Self	2024-present
Nistha Kumar	2024-present
Vidhata Singh	2024-present
Rami Wilson	2024-present
Khin (Sabrina) Yone	2024-present
Anish Kumar	2023-present
Gene Park	2024-present
Richard Ho	2024
Ryan Leung	2024
Michael Soebroto	2024
Yanwen Xu	2024
Jimmy Franknedy	2023-2024
Rohan Ghosalkar	2023-2024
Shiyuan Huang	2023-2024
Ameyaa Biwalkar	2022-2024
Dhanishtha Patil	2023-2024
Harshini Venkataraman	2023-2024
Anuj Kamat	2023-2024
Olivia Anastassov	2022-present
<i>On leave since June 2023</i>	
Kalyana Suma Sree Tholeti	2023-2024
Arka Pal	2023
Maithili Luktuke	2023
Supriya Paulose	2023
Ryan Ahari	2023
Shayan Salsabilian	2023
Joseph Silberman	2023
Sukhveer Karicut	2023
Vidyasri Ravi	2022-2023
John Yu	2022-2023
Rohit Das	2022-2023

	Smruthi Pobbathi	2022-2023
	Rangasri Chakravarthy	2022-2023
	Jacob Low	2022-2023
	Yash Chhabria	2022-2023
	Keely Wheelan	2022-2023
	Zhihang Zhou	2022
Project Reader	Nafisa Hussain	2024
	Jay Mehta	2024
	Valentina Tang	2024
	Zichao Li	2024
	Daniel Sabo	2024
	Najmeh Mashhadi	2023
	Golam Md Muktadir	2023
	Srilekha Vutukuru	2023
	Aaron Monajjemi	2023
	Christopher Liu	2023
	Alexander Miller	2023
	Sreevani Suvarna	2023
	Rahul Bharadwaj	2022
	Alec Siegel	2022
	Ian Hardy	2023
	Jinghao Shen	2023
	Undergraduate Students (UCSC)	
Primary Advisor	Rishika Srinivas	2022-present
	Arnav Kartikeya	2022-present
	Shreedhar Jangam	2023-present
	<i>2024 Dean's Award Winner.</i>	
	Siddarth Mamidanna	2023-present
	<i>2024 Dean's Award Winner.</i>	
	Neha Suresh	2024-present
	Rajat Maheshwari	2024-present
	Anshika Agarwal	2024-present
	Selam Mitike	2024-present
	Daniel Yang	2024-present
	Ashwin Pazhoor	2024-present
	Bill Zhang	2024-present

Jihang Li	2024-present
Anish Talluri	2024-present
Ian Yam	2024-present
Thomas Wang	2024-present
Umair Rizwan	2024-present
Sri Seethammagari	2024-present
Samiyah Shaikh	2024-present
Macklin Reeve-Wilson	2024-present
Sebastian Woo	2024-present
Sam Cheng	2023-present
Ish Khandelwal	2023-present
Tommy Chen	2022-present
Raj Singh	2023-2024
Coen Adler	2021-2024
<i>Bachelor's thesis student. 2024 Chancellor's Award Winner.</i>	
Batuhan Salih	2022-2024
Nick Wang	2022-2024
Tong Jia	2023-2024
Roger Li	2023-2024
Camden Beard	2023-2024
Daniel Hsieh	2023-2024
Max Myers	2023-2024
Oliver Lokhandwala	2023-2024
<i>Koret Scholarship Recipient.</i>	
Fariha Lateef	2023-2024
Andrew Susanto	2023-2024
Brandon Shapiro	2023-2024
Daniel Sarni	2024
Vaibhav Honakere	2024
Erick Hernandez	2024
Janvi Rochlani	2024
Sebastian Osorio	2024
Sahil Gupta	2022-2023
<i>Bachelor's thesis student.</i>	
Alex Fegghi	2023
Michael McAllister	2023

	Suneet Katrekar	2022-2023
	Owen Shi	2022-2023
	Anish Kumar	2023
	Tanisha Khemka	2022
	Ariel Kamen	2022
	Lan Mi	2022
	Radhika Gadre	2022
	Aylin Akkus (visiting student)	2022
	Christopher Oey	2022
	Aaja Ouellette	2022
	Isaac Plotkin	2022
Reading Committee	Nishanth Jayram	2022

MIT Thesis Students (12+ month fulltime student)

MEng	Tianye Chen	2018-2019
	<i>Co-advised with Lalana Kagal. Co-authored paper on rule-learning [34].</i>	
SuperUROP	Evelyn Florentine	2017-2018
	<i>Co-authored journal paper on monitoring opaque learning systems [7].</i>	
	Zoe Lu	2017-2018

MIT Research Project Students (6+ month semester course)

UROP	Shayda Moezzi	Fall 2020-2021
	Vishnu S Penubarthi	Fall 2019-2021
	<i>Co-authored paper on multimodal explanations [17].</i>	
	Marla E. Odell	Spring 2019
	Elizabeth Han	Spring 2019
	Obada Alkhatib	IAP/Spring 2018
	Michal Reda	IAP/Spring 2018
	Ishan Pakuwal	IAP/Spring 2018
UAP	Matthew Kalinowski	Spring 2017

University Service

University Service	Senate Committee on Rules, Jurisdiction and Elections (RJE)	2023-present
Dept. Service	PhD Visit Days/Open House Organizer	2023
	MS/PhD Recruitment Committee	2023-present
	NLP MS Capstone Project Mentor	2023-present

	Graduate Admissions Committee	2022-present
	Security and Privacy Faculty Search Committee	2023-present
	LSOE Faculty Search Committee	2023-present
	LSOE Faculty Search Committee	2022-2023
	Grad AI Day Panel Organizer	2024
Committee	External Member: HCI Associate Prof. Search (UCSC CM Dept.)	2022-2023

Professional Activities

Editorial Board	Journal of Neurosymbolic Artificial Intelligence	2022-present
Guest Editor	Special Issue on Neurosymbolic AI for Cyberphysical systems	2023-present
	AI Magazine Special Issue on Anticipatory Thinking	2022-2023
Program Chair	Advances in Cognitive Systems (ACS) [link]	2022
Organizer	AAAI workshop on Explainable AI for Deep RL: [link]	2024
	NeurIPS workshop on Building Trustworthy Embodied AI: [link]	2022
	AAAI Fall Symposium on Anticipatory Thinking [link]	2020/2021
	NeurIPS Workshop on XAI for debugging and diagnosis [link] .	2021
	ACS Workshop on Story Enabled Intelligence. [link] .	2019
	AAAI Spring Symposium 2019: Story-Enabled Intelligence. [link] .	2019
	MIT Machine Learning Interpretability Reading Group	2018-2020
	MIT AI and Ethics Reading Group. [link] .	2018-2020
	MIT IPRI Privacy, Security and Policy (PSP) Meeting	2018-2019
	MIT Path of Professorship Workshop	2018
	MIT EECS Visit Days and Orientation	2016
Advisory Board	Center for Advancing Safety of Machine Intelligence (CASMI) <i>Governance Advisory Committee</i> . [CASMI Leadership]	2022-present
	Bias Advisory Board (EU) [Bias Advisor Board]	2024-present
Local Chair	NeuroSymbolic AI Conference (NeSy)	2025
	Advances in Cognitive Systems	2019
Senior PC	AAAI	2022-present
PC	Advances in Cognitive Systems (ACS)	2022-present
	AIC	2023
	Facct	2023
	Scheme Workshop	2023
	Semantic Web Journal issue on Commonsense Knowledge and Reasoning	2021
	AAAI	2021-2022

	ACS Workshop on Goal Reasoning	2021
	AAMAS	2021
	DX-2020	2020
Grant Reviewer	AFOSR Trust and Influence program	2023
	NSF CISE FRR Panel Reviewer	2022
Paper Reviewer	The Web Conf	2023
	AAAI SRRAI track	2023
	EAAI	2023
	NeurIPS Workshop	2023
	AI4HRI	2022
	Information Issue on Foundations and Challenges of Interpretable ML	2021
	HRI	2020-present
	ACS Workshop on Goal Reasoning	2021
	Artificial Intelligence Review	2020
	IEEE Transactions on Cybernetics	2019
	NeurIPS	2019
	AAAI Spring Symposium	2019
	Slovak-Israeli Scientific Research Program	2018
	MIT MITES	2018
	HRI Late Breaking Reports (LBR)	2018
	AAAI (Guest Reviewer)	2015
Student Rep.	Stanford ICME	2011-2013
	MIT EECS Visiting Committee	2017
	<i>Met with the EECS Visiting Committee and gave a personal perspective on the EECS Department, student life, and diversity.</i>	
	MIT Grad Rat	2017-2019
Volunteer	UCSD Alumni Board	2015-2019

Outreach

UCSC Fall Open House Faculty Panel	2024
Invited Talk at Los Altos High AI Club	2024
Invited Plenary Speaker-BayHacks	2024
Rotary Club of Salinas talk	2024
Rotaract Club of Santa Cruz County: Judging Panel	2024
UCSC Women in Science Society Panel	2024
UCSC INSPIRE AI Panel	2024

ACM Club at UCSC: "What Do Profs Do All Day?"	2023
SC Works AI Horizons Part 2 Speaker	2023
UCSC ACM/GraceHacks Keynote Speaker	2023
UCSD CSE Ph.D. Industrial Job Search Panel	2023
Tri-Valley Hacks Judge	2023
Santa Cruz County Science & Engineering Fair Judge	2022

Advocacy

I am passionate about AI and policy, AI and ethics, and informing legislation pertaining to liability for Autonomous Vehicles (AVs).

UC congressional briefing on artificial intelligence 2023

Joint with Ken Goldberg (UC Berkeley), William Wang (UCSB), David Danks (UCSD), and Brandie Nonnecke (UC Berkeley).

Diversity, Equity and Inclusion

UCSC MESA Faculty Forum steering committee 2022-present

Joint with Imelda Meza, Academic Advisor/Program Coordinator and Marcella Gomez, Associate Professor, AM

Almanac of Technology and Advancement (ATA) Faculty Sponsor 2024-present

UCSC ACM-W Board Faculty Sponsor 2022-present

Mentor Women in ML (WiML) PhD Mentor 2022-present

Black in AI 2020

MIT EECS GAAP 2020

The Graduate Application Assistance Program (GAAP) is a student-run, volunteer-based program which provides assistance to EECS PhD applicants from under-represented groups, including students from groups underrepresented in STEM and students with non-traditional academic backgrounds.

Xerox ABI Mentoring Program 2014-2015